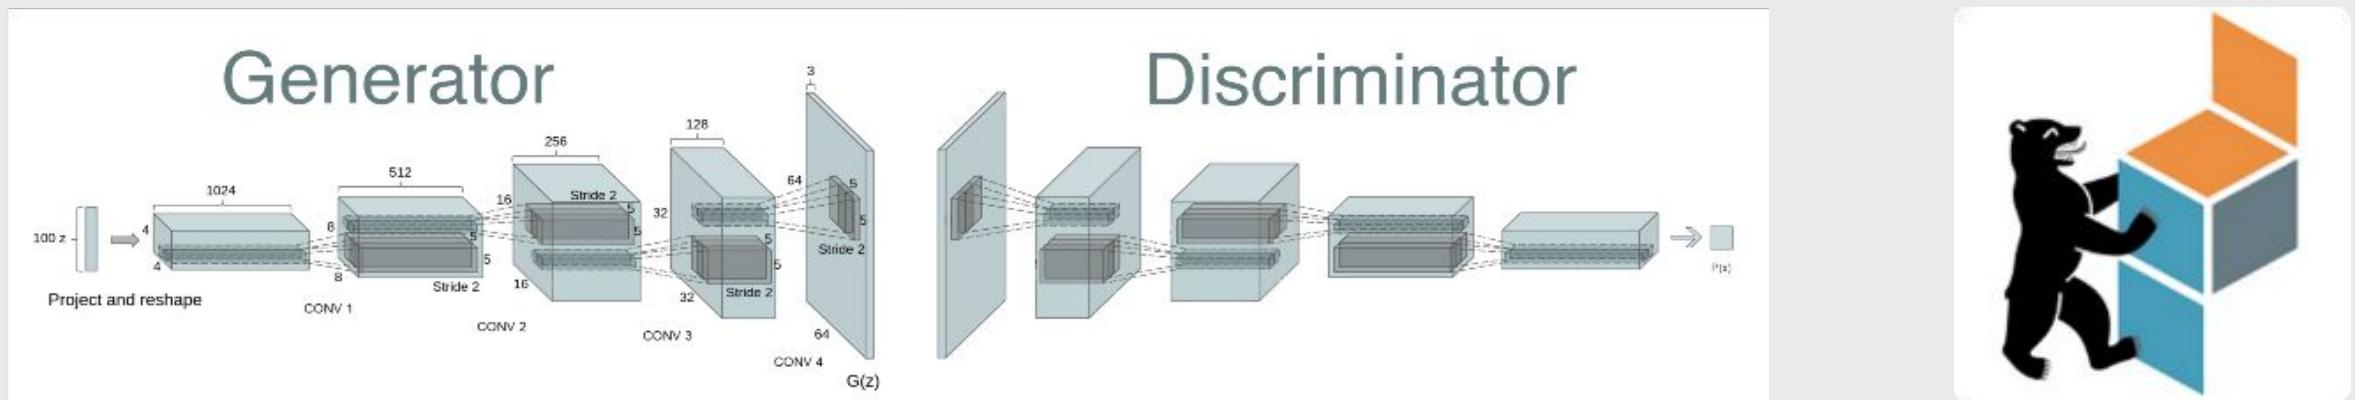


[PyData Berlin October meetup]

Deep Convolutional Network and Deep Generative Networks
to Automate Industrial Manufacturing Applications



Slides available at: https://docs.google.com/presentation/d/1pNif34jNv_RUH2caHDVvhVzJGhfW_5QL6DIuQEJDjhc/

about:myself – Jonathan Dekhtiar

26 years old

I am French Food/Wine lover

Data Scientist Engineer

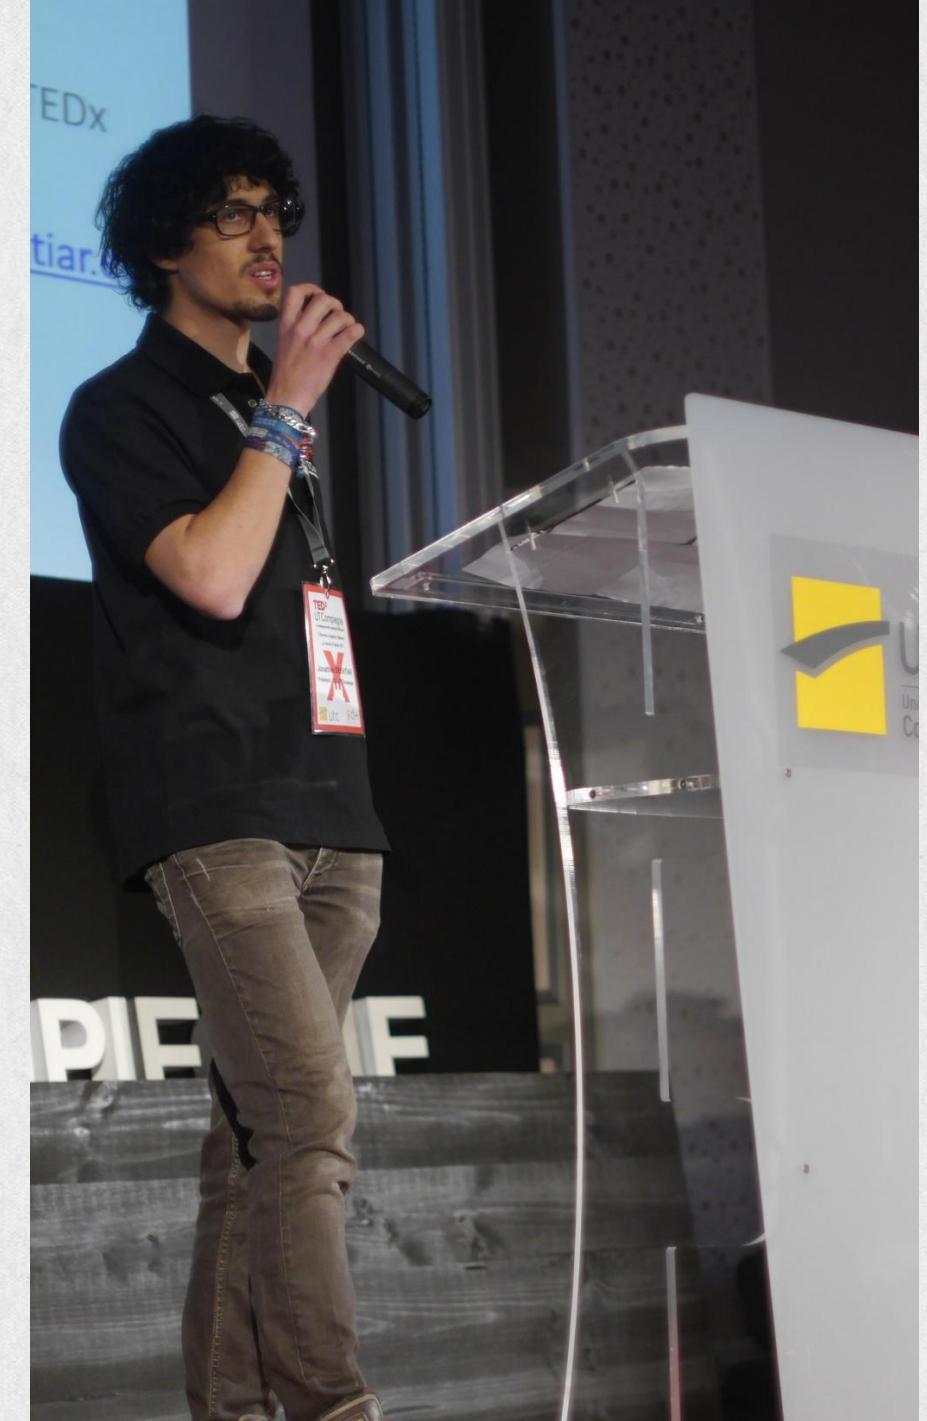
PhD Student in Deep Learning applications - UTC (FR) – EPFL (CH)

I try to automate & apply DL technologies in the manufacturing field.

Twitter: [@Born2Data](#)

Github: <https://www.github.io/DEKHTIARJonathan>

Email: contact@jonathandekhtiar.eu



What's on the menu today ?

1. Quick introduction: Data Science & Manufacturing Industry
2. Supervised Learning and an open dataset targeted manufacturing industry
3. Unsupervised Learning and Generative Adversarial Learning
4. Why and How could you use a GAN to detect novelty/defects ?
5. A few tips n' tricks to fine-tune a GAN

1. Data Science and Manufacturing Industry – Easy fit ?

- A few major challenges -

Data Heterogeneity and Data Complexity

- ** A wide variety of (proprietary) formats.
- ** Engineering files are way more complex than Images/Videos/Sound.

Difficulties to collect large volumes of data

- ** The industry have very complex data in small quantities (opposite of GAFAM)
- ** Collecting just a little data can cost a lot of money.
- ** Intellectual Property issues: let's negotiate!

Reduced System Inter-Operability

System were not designed to communicate outside of closed proprietary environment.

Very Expensive Licenses

You need to pay for very \$\$\$ developers licenses (\$30k for one developer license in C++ with Catia)

2. Supervised Learning and opening a CV dataset

Supervised, Unsupervised or Reinforcement Learning ?

**** SL is often not possible.**

- ⇒ domain-experts are paid a lot of €€€
- ⇒ domain-experts are hardly available

**** RL approaches are very interesting**

- ⇒ take into account the domain-expert feedback
- ⇒ model fine-tuning over time

**** UL approaches present some appealing advantages**

- ⇒ does not require an expensive expert
- ⇒ cheaper
- ⇒ easier to develop and put into production

2. Supervised Learning and opening a CV dataset

Supervised, Unsupervised or Reinforcement Learning ?

**** SL is often not possible.**

- ⇒ domain-experts are paid a lot of €€€
- ⇒ domain-experts are hardly available



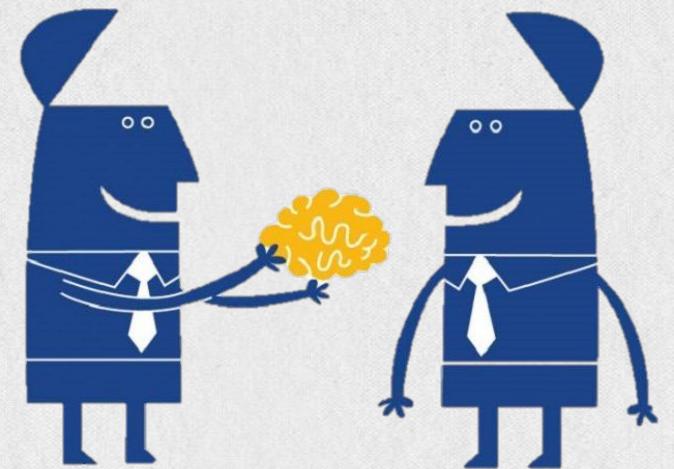
Transfer Learning will be
your answer to many situation

**** RL approaches are very interesting**

- ⇒ take into account the domain-expert feedback
- ⇒ model fine-tuning over time

**** UL approaches present some appealing advantages**

- ⇒ does not require an expensive expert
- ⇒ cheaper
- ⇒ easier to develop and put into production



2. Supervised Learning and opening a CV dataset

DMU-Net – Open Dataset for the Manufacturing Industry

- <https://www.dmu-net.org> -

Actuator → 37	AirFilter → 32	Battery → 46	Bearing → 128	Brake → 39
Camshaft → 46	Clutch → 61	Cooler → 30	Coupling → 43	Crankshaft → 81
CuttingTool → 41	ElectricMotor → 142	ElectronicBoard → 37	ExhaustPipe → 20	Gear → 18
GearWheel → 91	Helix → 60	Nut → 56	Piston → 87	Pulley → 109
Pump → 34	Rod → 53	RodPistonAssembly → 56	Screw → 219	ShockAbsorber → 78
SparkPlug → 87	Switch → 87	Turbo → 18	Washer → 47	WingNut → 33

+/- 2000 CAD Models in STEP-AP242, JT, STL, ThreeJS, X3DOM and Object File (OBJ+MTL)

2. Supervised Learning and opening a CV dataset

DMU-Net – Open Dataset for the Manufacturing Industry

- <https://www.dmu-net.org> -



2. Supervised Learning and opening a CV dataset

DMU-Net – Open Dataset for the Manufacturing Industry

- <https://www.dmu-net.org> -

3D Model Visualiser

The screenshot shows a 3D model of a brake assembly in the center, with a list of sub-assemblies on the left and component icons on the right.

Hide/Show Sub-Assemblies:

- BRAKE_HUB-1_XT_0
- DRIVER_BRAKE_ROTOR-1_XT_0
- BRAKE_CALIPER_OUTER-1_XT_0
- RETAINING_PIN-1_XT_0
- RETAINING_PIN-1_XT_0
- BRAKE_PISTON-1_XT_0
- LEFT_BRAKE_PAD-1_XT_0
- RIGHT_BRAKE_PAD-1_XT_0
- BRAKE_PAD_RETAINER_1-1_XT_0
- BRAKE_PAD_RETAINER_2-1_XT_0
- BRAKE_PISTON-1_XT_0
- BRAKE_CALIPER_OUTER-1_XT_0
- BRAKE_PISTON-1_XT_0
- BRAKE_PISTON-1_XT_0

Component Icons (Left and Right):

- ElectricMotor
- Actuator
- GearWheel
- ElectricMotor
- Cooler
- Screw
- ExhaustPipe
- Actuator
- Turbine
- ElectricMotor
- Gear
- ElectricMotor
- ExhaustPipe
- Bearing
- GearWheel
- ExhaustPiping
- Crankshaft
- Helix

Jonathan DEKHTIAR - @Born2Data

3. Unsupervised Learning & Generative Adversarial Nets

Supervised, Unsupervised or Reinforcement Learning ?

**** SL is often not possible.**

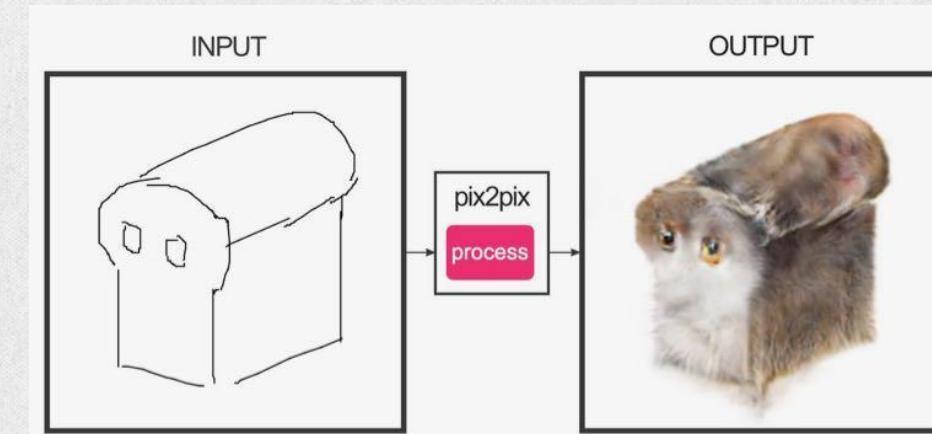
- ⇒ domain-experts are paid a lot of €€€
- ⇒ domain-experts are hardly available

**** RL approaches are very interesting**

- ⇒ take into account the domain-expert feedback
- ⇒ model fine-tuning over time

**** UL approaches present some appealing advantages**

- ⇒ does not require an expensive expert
- ⇒ cheaper
- ⇒ easier to develop and put into production



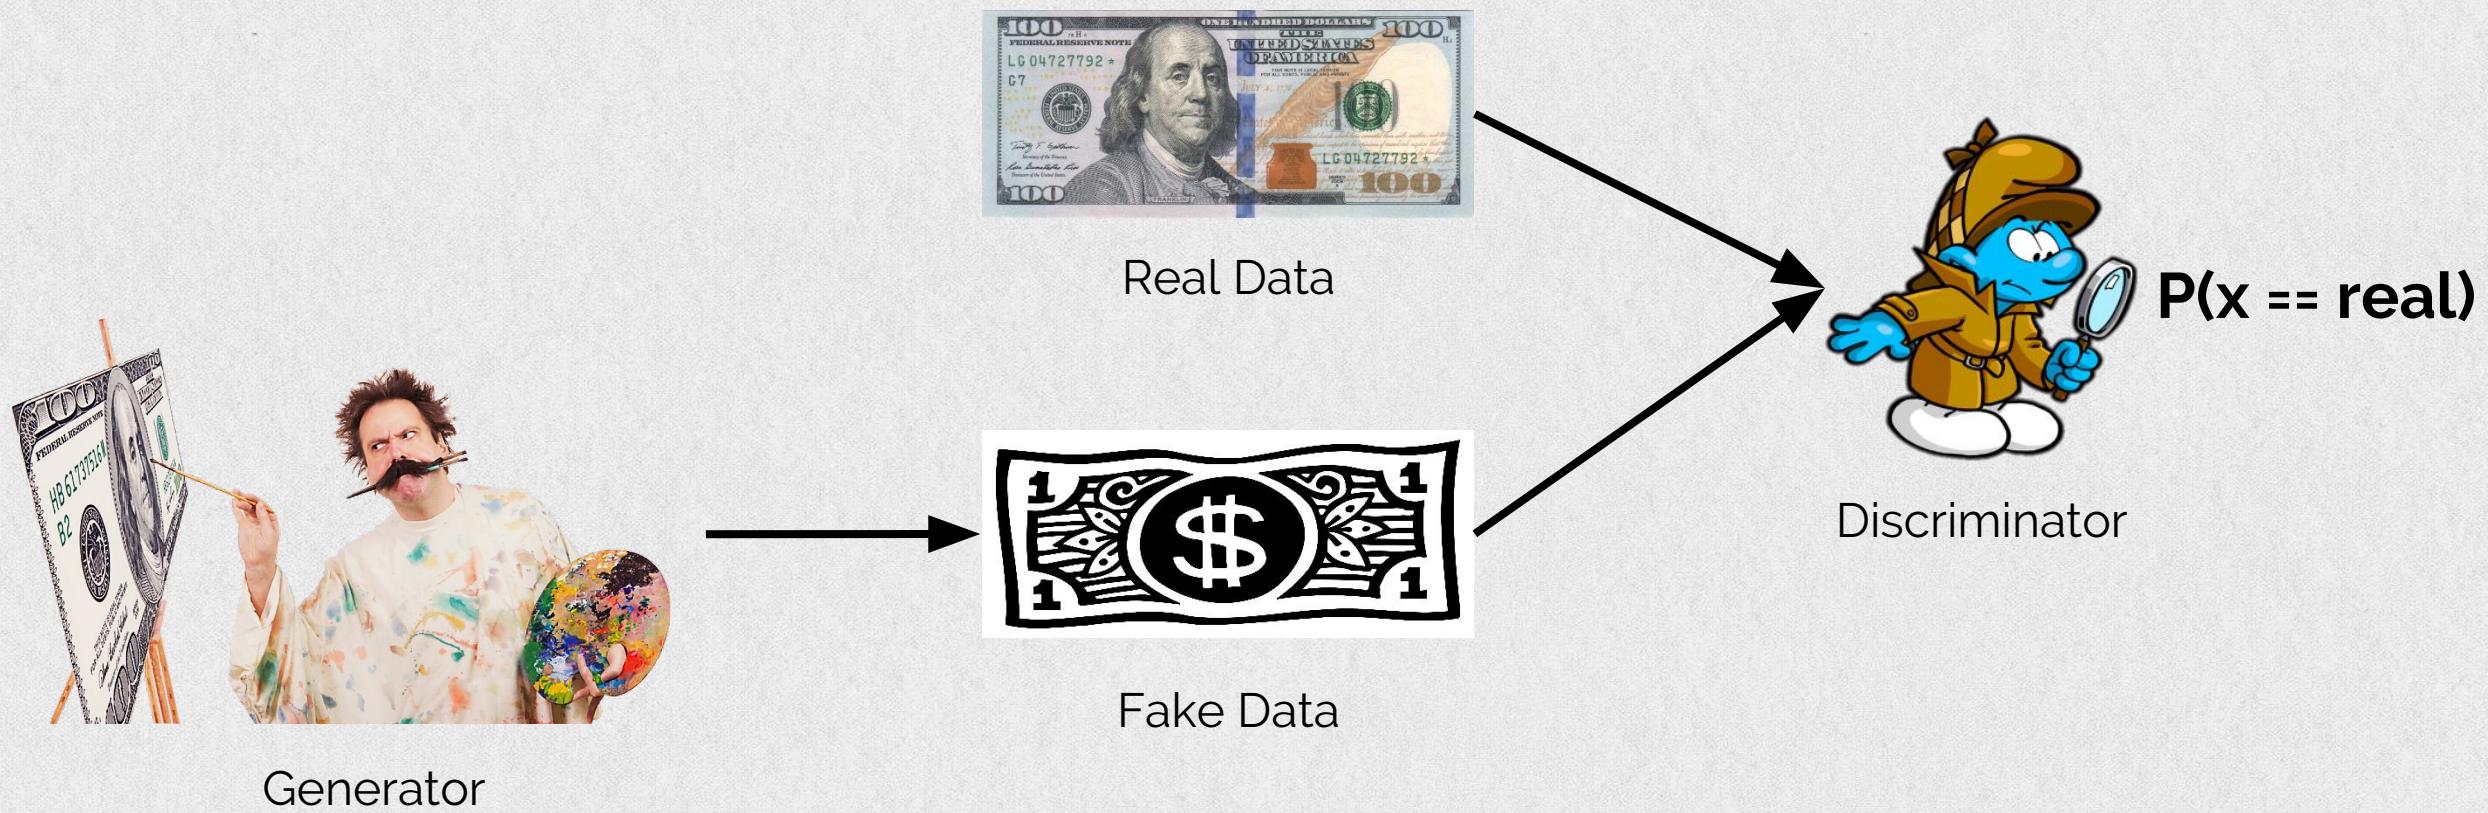
Generative Adversarial Neural Networks
are one of the possible directions.

Let's explore this possibility



3. Unsupervised Learning & Generative Adversarial Nets

Generative Adversarial Neural Networks



3. Unsupervised Learning & Generative Adversarial Nets

Minimax Game and Adversarial Training

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$

The Generator G and the Discriminator D have opposite objectives:

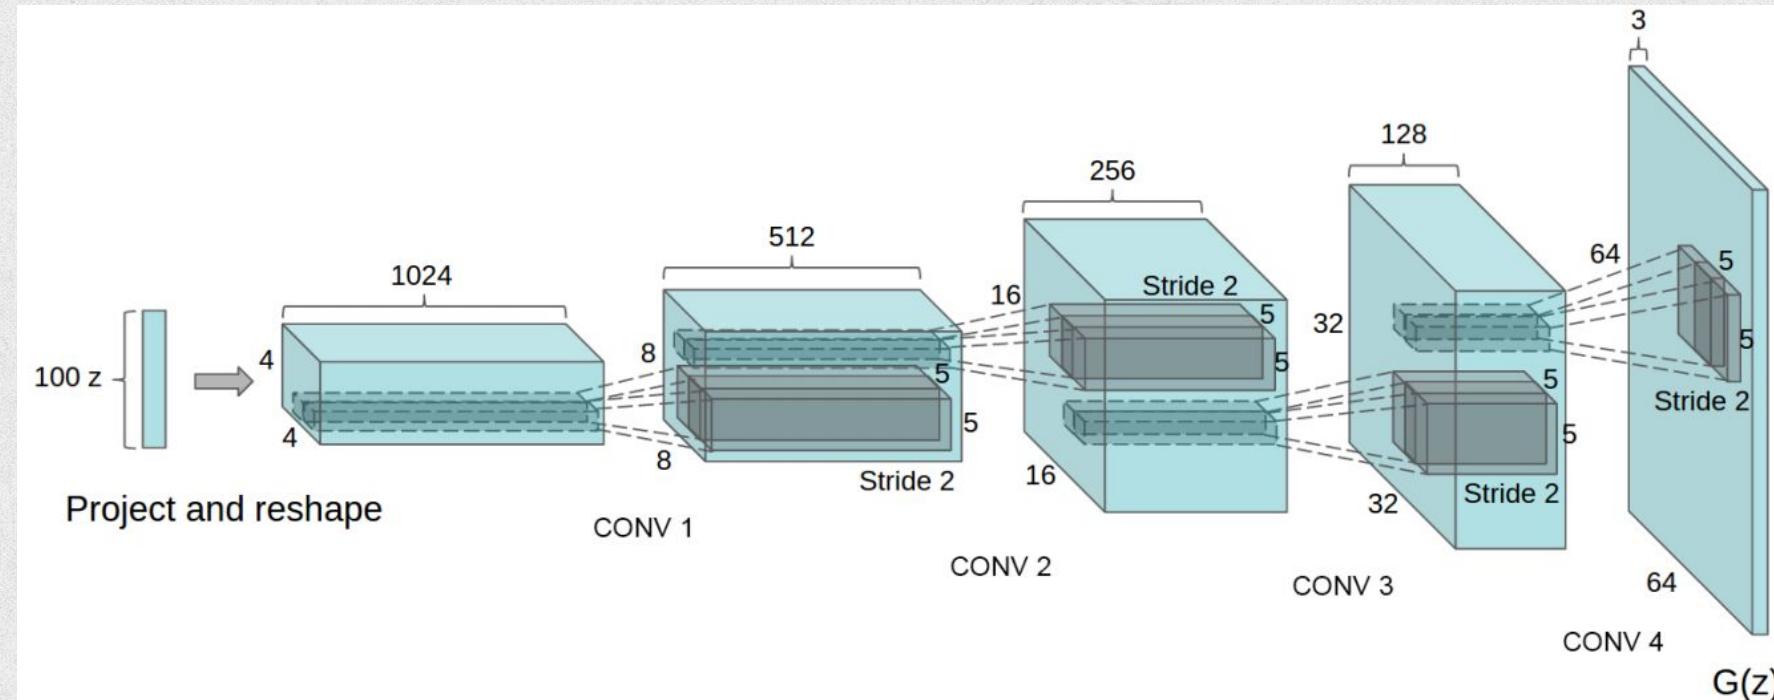
- ⇒ If G wins then D loses and vice versa.
- ⇒ Each network tries to maximise his log likelihood.

When the training converges, this point is called a **Nash Equilibrium**:

- ⇒ D has 50% chance of getting its decision right (*fake or real data*).
- ⇒ In other words, data produced are +/- indistinguishable from real data.

3. Unsupervised Learning & Generative Adversarial Nets

GAN and DCGAN

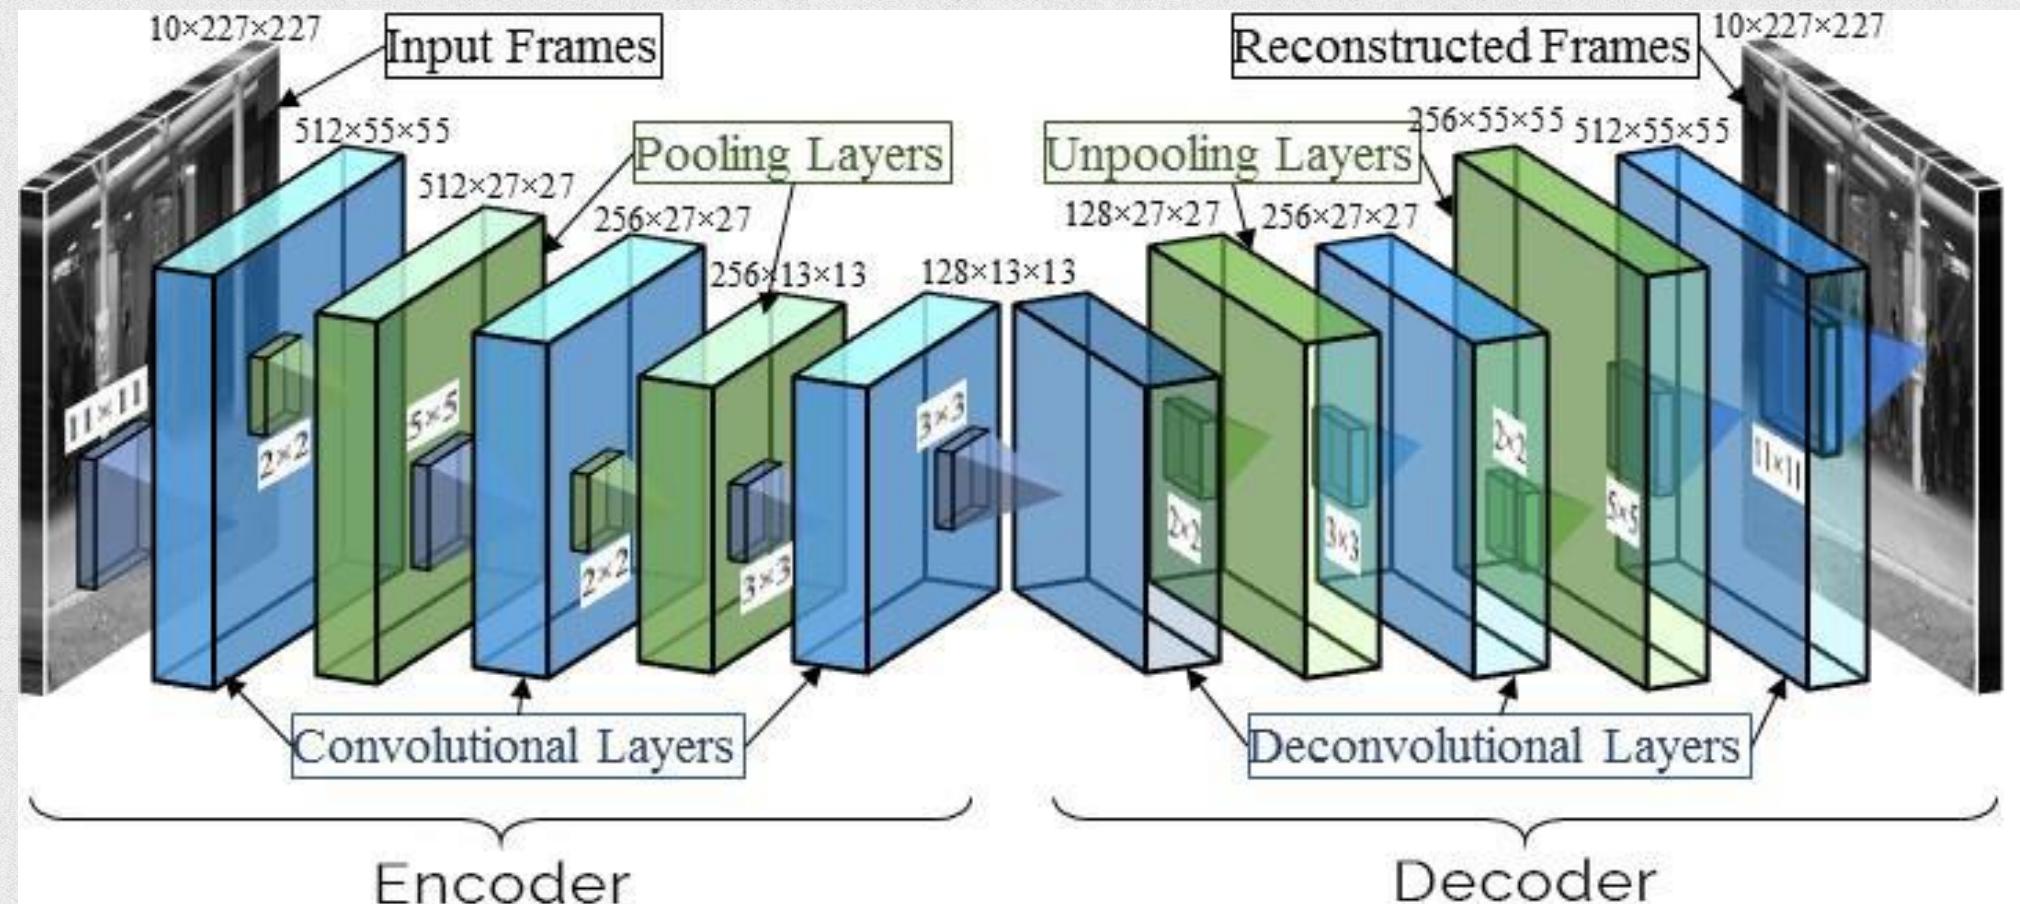


Generative Adversarial Networks – Goodfellow et al. 2014

Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks – Radford et al. 2015

3. Unsupervised Learning & Generative Adversarial Nets

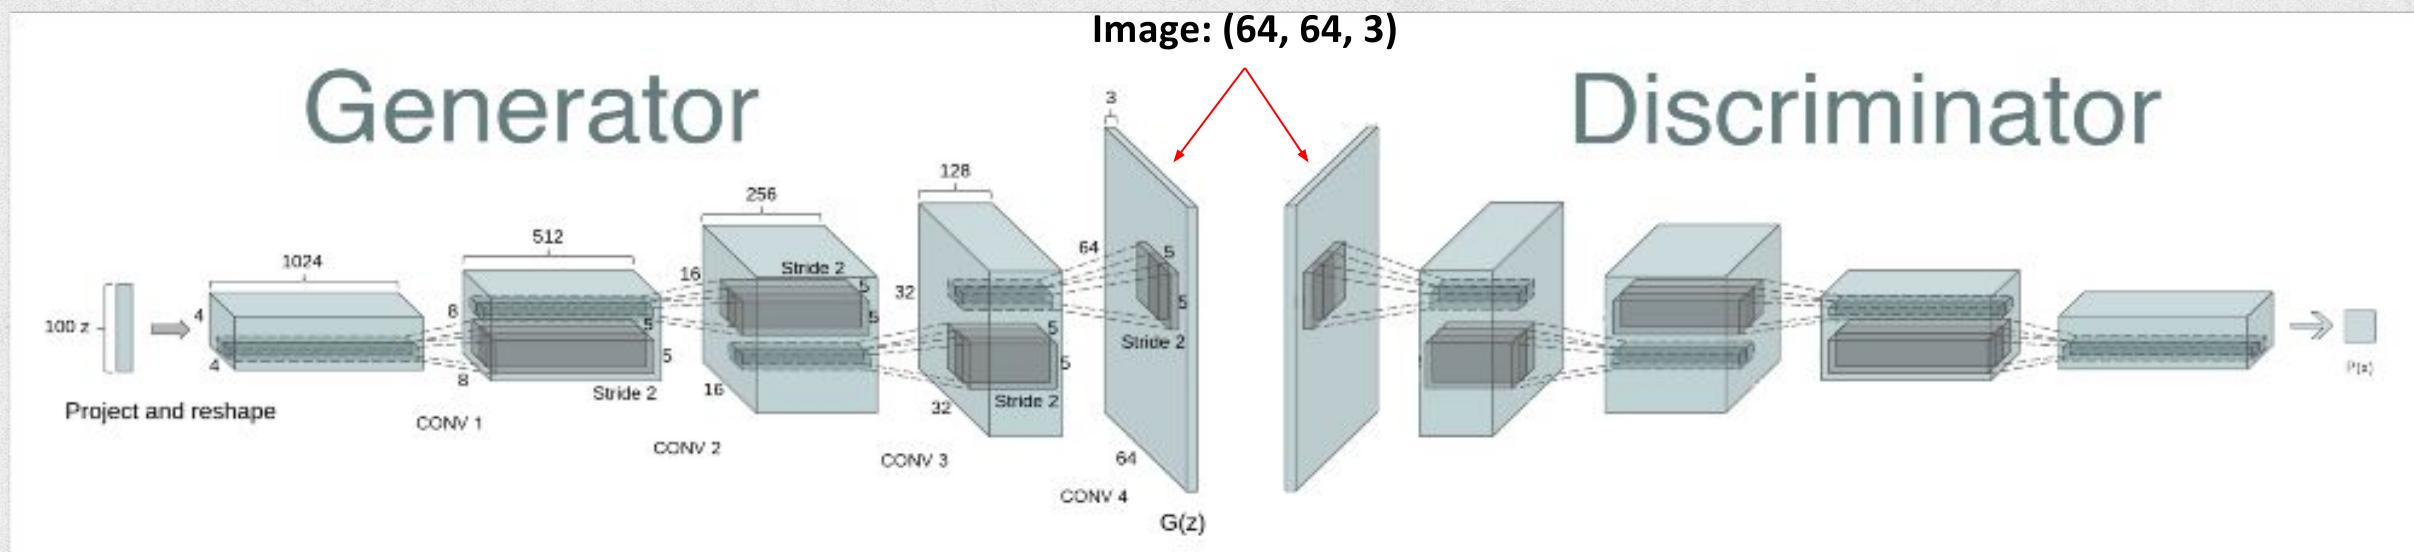
Convolution vs Strided-Convolutions – Example with an AE



Source: Learning Temporal Regularity in Video Sequences – Hasan et al. 2016

3. Unsupervised Learning & Generative Adversarial Nets

DCGAN Implementation



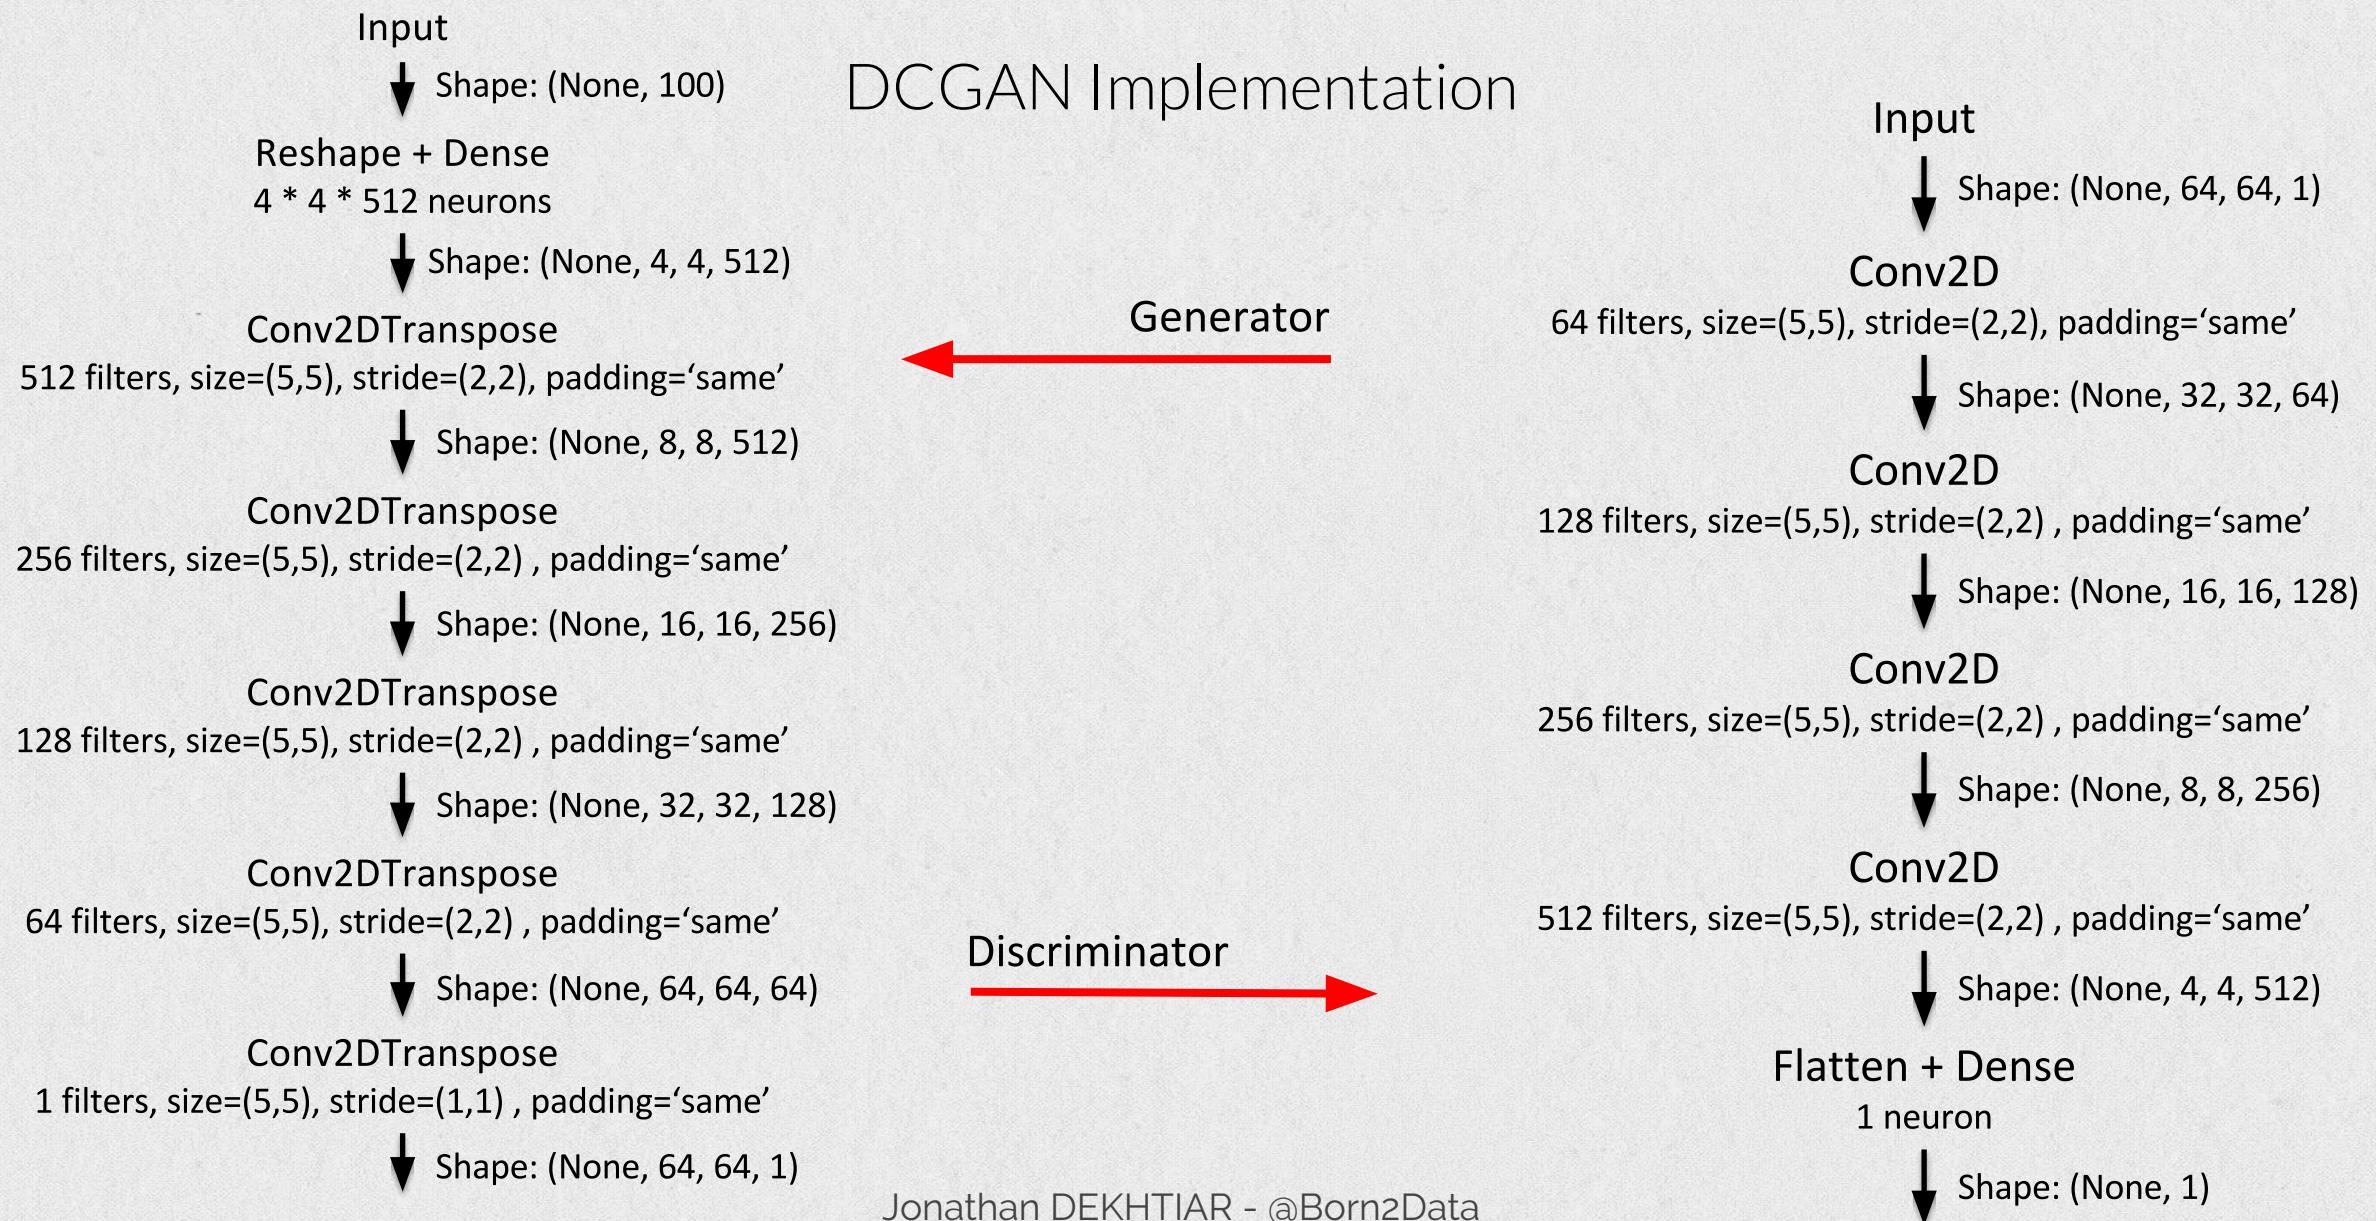
Network Settings

Stride	= (2, 2)
Kernel Size	= (5, 5)
Padding	= "same"
Activation	= Leaky ReLU
N_Filters	= [1024, 512, 256, 128]
Latent vector shape	= (None, 100)

Training Settings

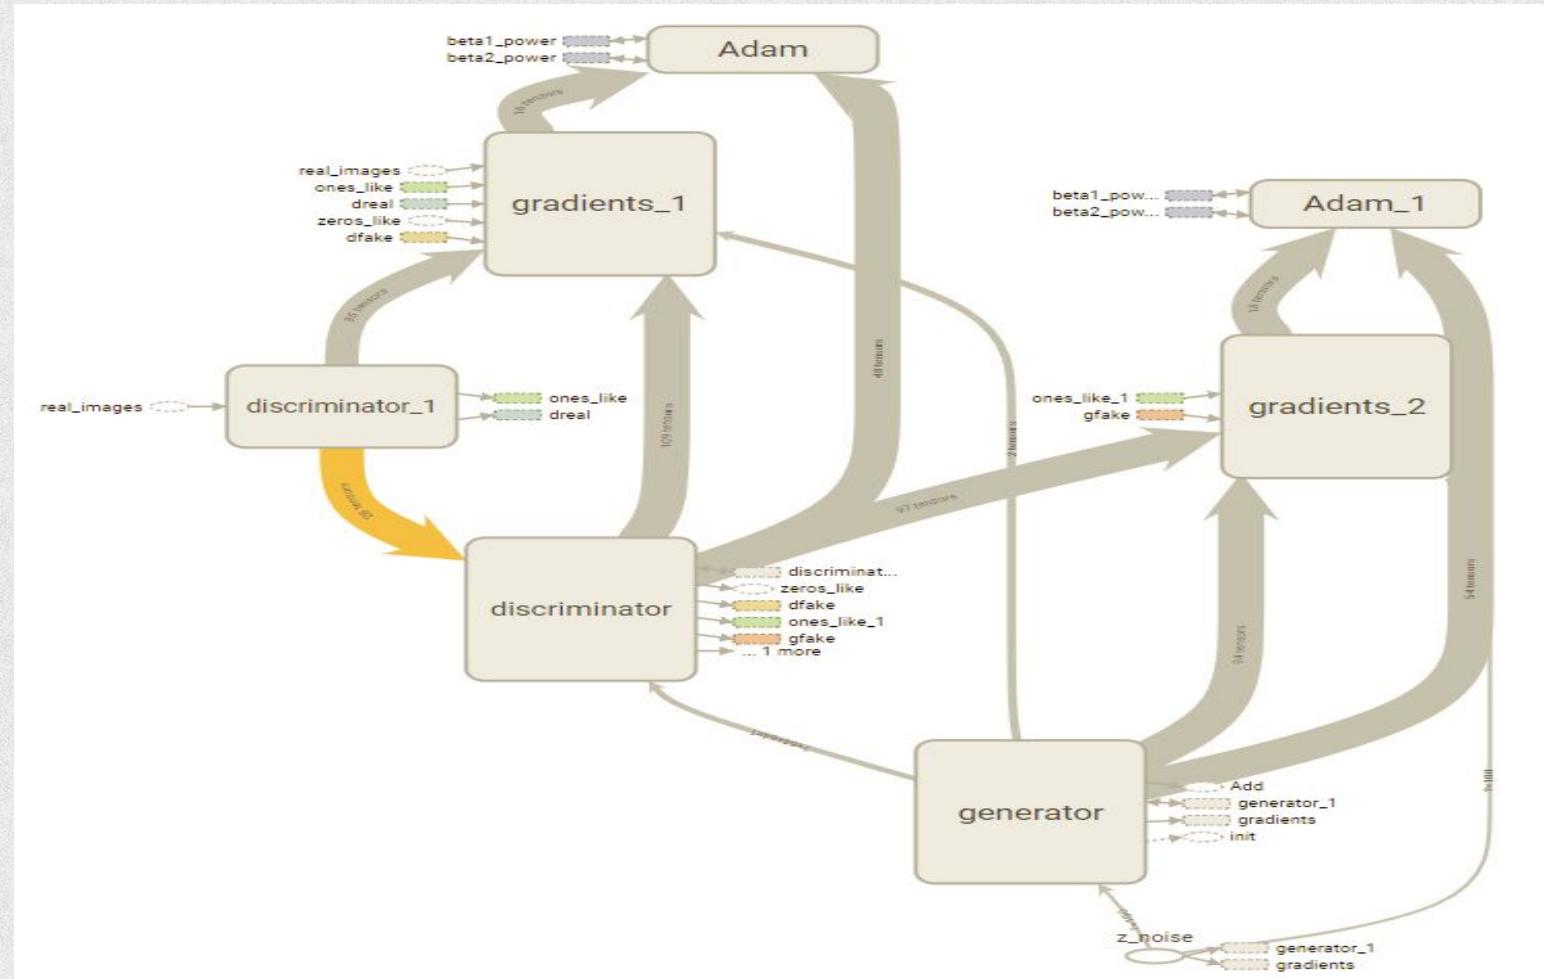
Optimizer	= Adam
Learning Rate	= 2e-4
Beta1	= 0.5
Beta2	= 0.999

3. Unsupervised Learning & Generative Adversarial Nets



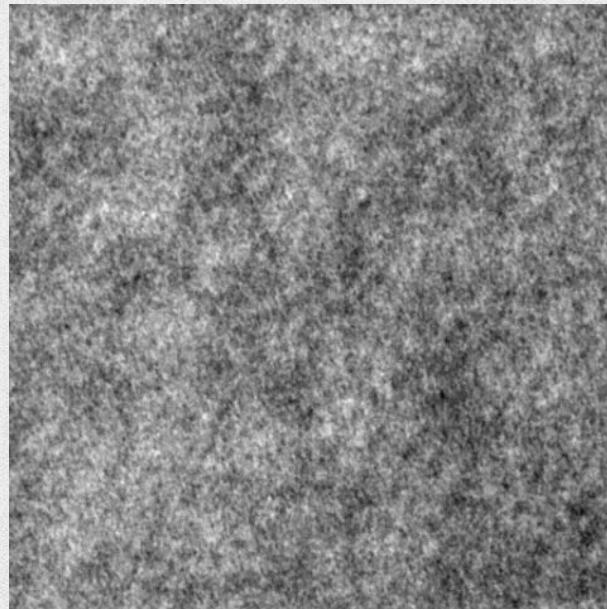
3. Unsupervised Learning & Generative Adversarial Nets

TensorBoard view of DCGAN

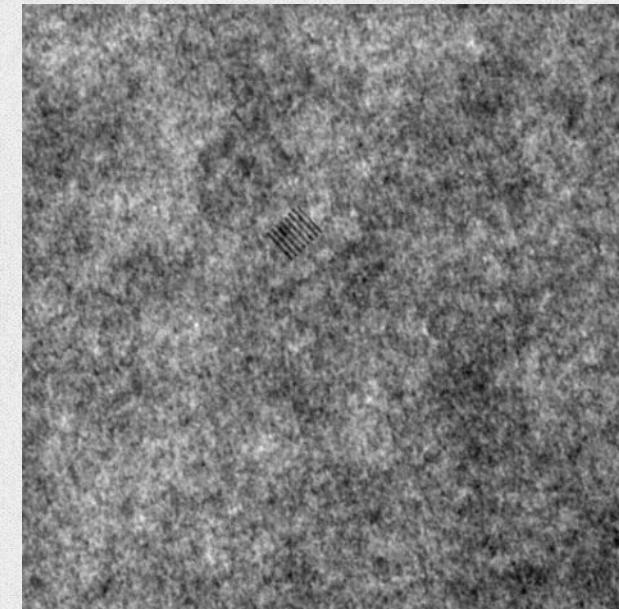


4. How did I use GANs to detect novelty/defects ?

Dataset used for this experiment



Healthy Machined Surface



Defective Machined Surface

Toy Dataset Provided by:

German Association for Pattern Recognition <http://resources.mpi-inf.mpg.de/conferences/dagm/2007/>

4. How did I use GANs to detect novelty/defects ?

How GANs can help detecting defects / novelty / unusual situations

1. Make G learns how to generate the best looking “healthy” image.

2. Give a query image (healthy or defective) and try to re-create it

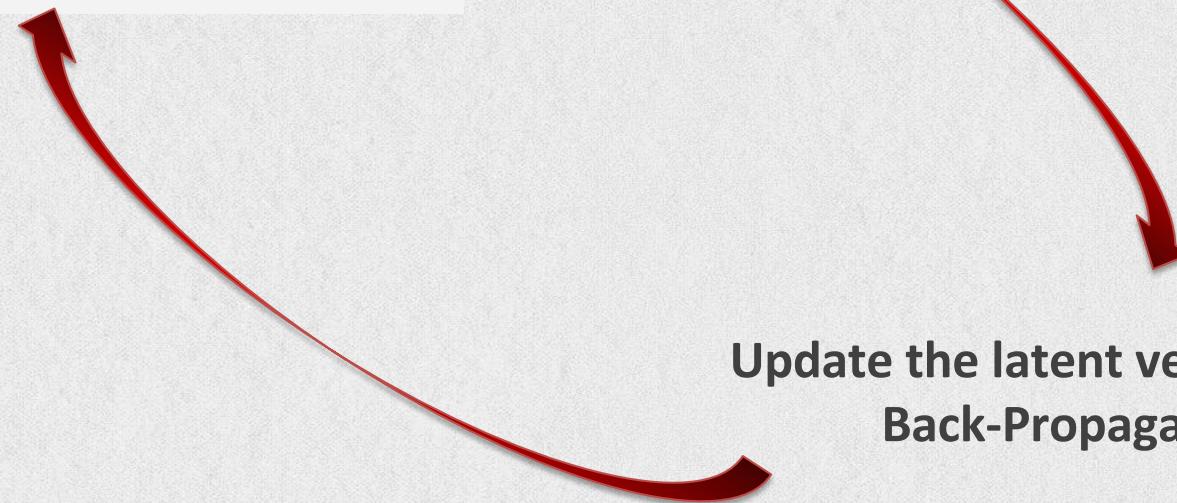
3. Can G produce a credible fake version ?

- ⇒ Yes: Query Image is healthy
- ⇒ No: Query Image is defective

4. How did I use GANs to detect novelty/defects ?

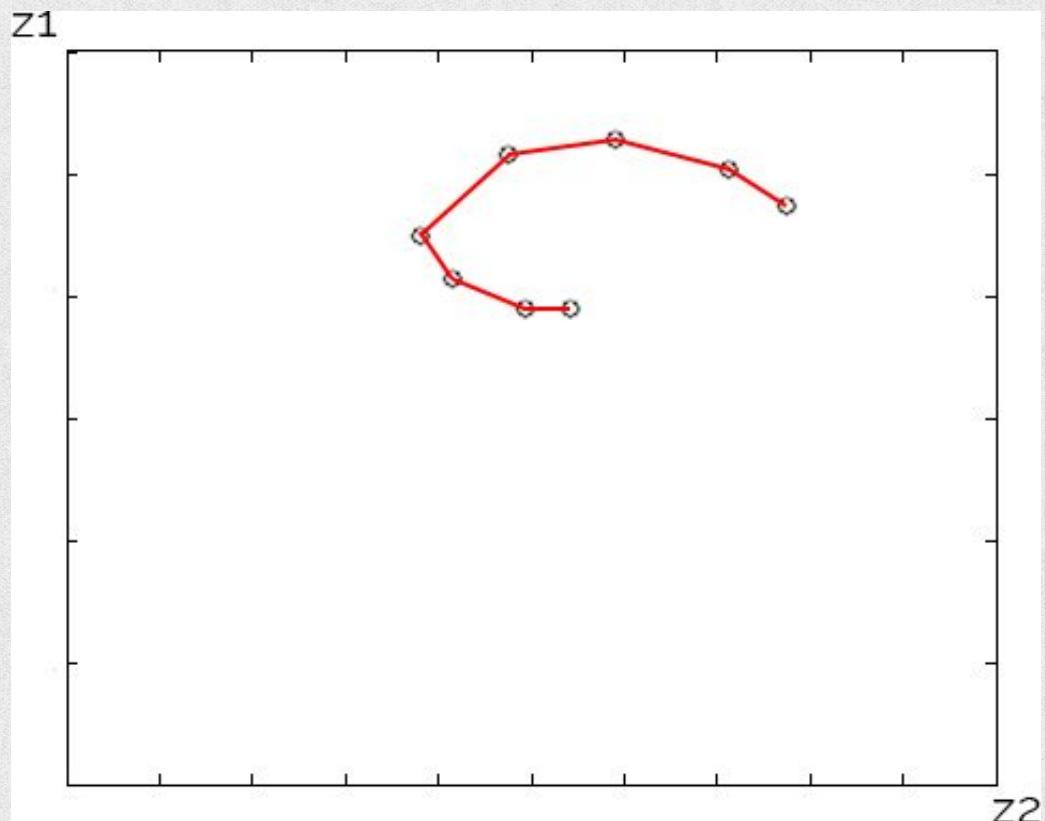
How can I control my GAN to produce the closest image as possible ?

```
def reconstruction_loss(query, target):  
    # Sum over pix2pix abs difference  
    loss = tf.subtract(query, target)  
    loss = tf.abs(loss)  
    loss = tf.reduce_sum(loss)  
    return loss
```



4. How did I use GANs to detect novelty/defects ?

Back-Propagate over Z during ... Evaluation



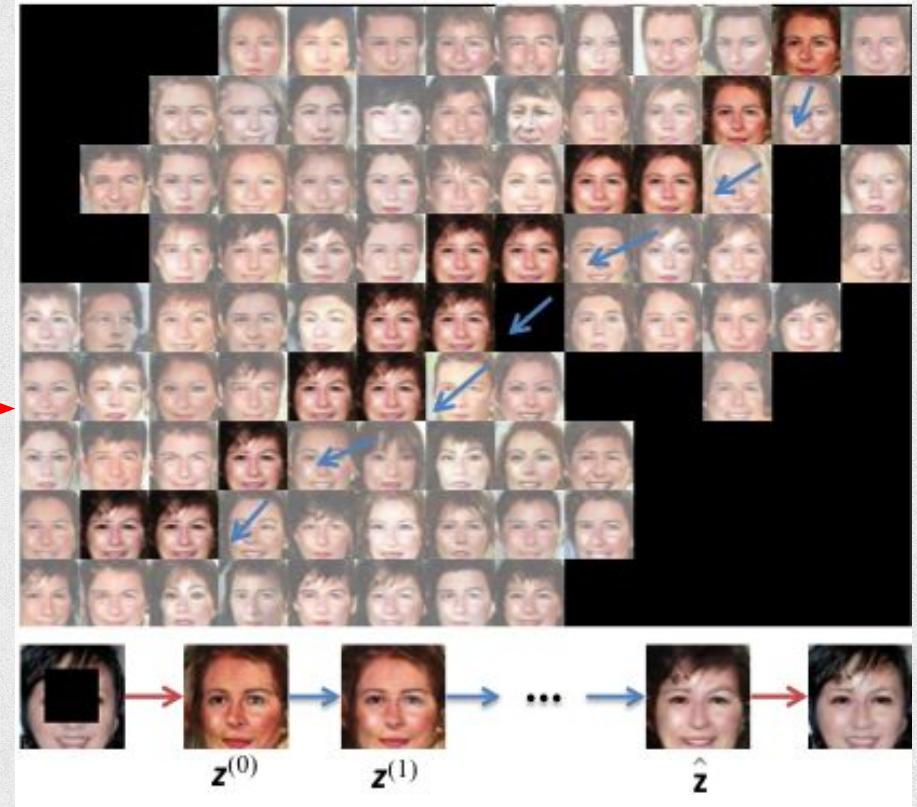
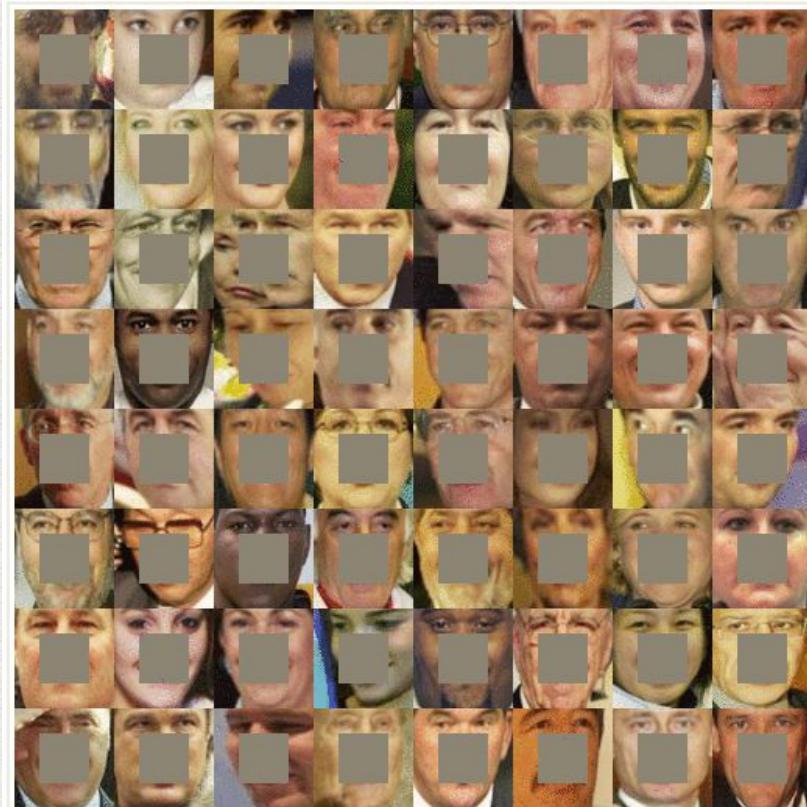
**** Perform Back-Propagation during evaluation:**

- ⇒ Try to generate the closest looking image by back-propagating through the input layer: Z.
- ⇒ Depending on the final loss value:
 - ** Low Value: Healthy*
 - ** High Value: Defective*

Inspired By: Semantic Image Inpainting with Deep Generative Models – Yeh et al. 2016

4. How did I use GANs to detect novelty/defects ?

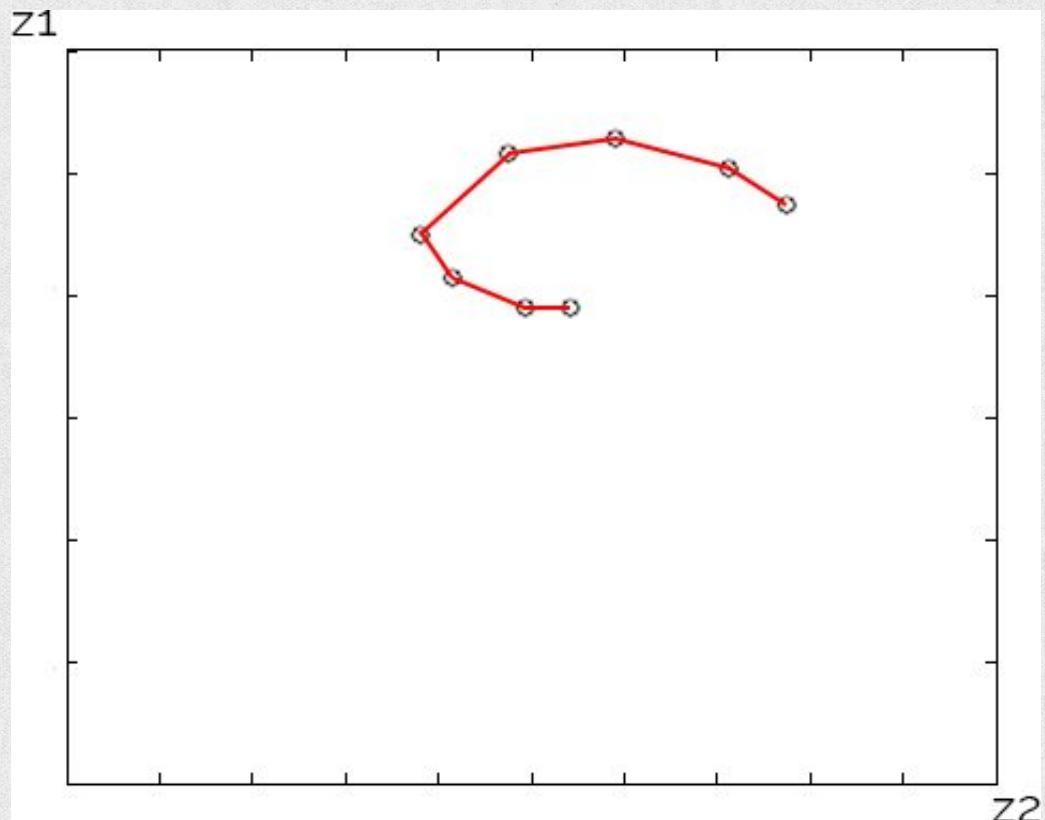
Back-Propagate over Z during ... Evaluation



Inspired By: Semantic Image Inpainting with Deep Generative Models – Yeh et al. 2016

4. How did I use GANs to detect novelty/defects ?

Back-Propagate over Z during ... Evaluation



**** Back-Propagation Settings:**

- ⇒ **Steps:** 500
- ⇒ **Learning Rate:** 1e-4
- ⇒ **Optimizer:** Momentum
- ⇒ **Gradient Momentum:** 0.8

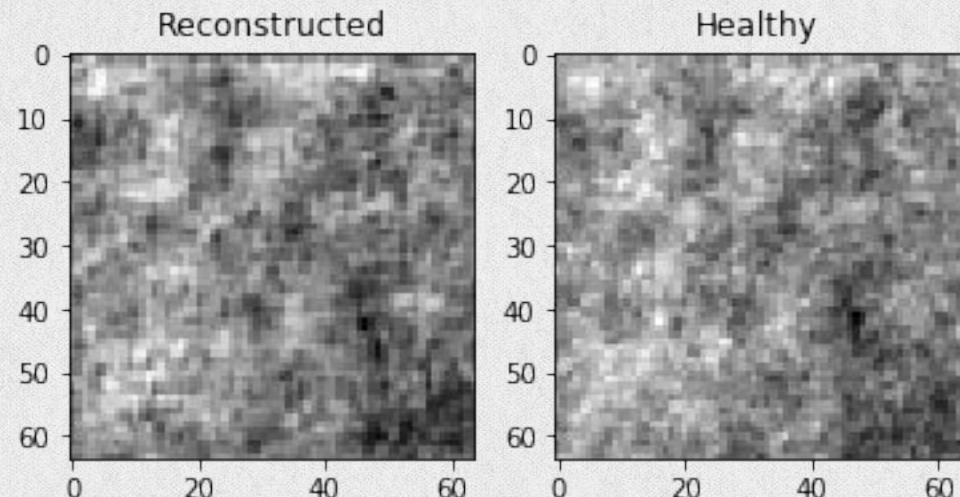
- ⇒ **Execution Time:** 1.5 sec per sample on GTX Titan X

Inspired By: Semantic Image Inpainting with Deep Generative Models – Yeh et al. 2016

4. How did I use GANs to detect novelty/defects ?

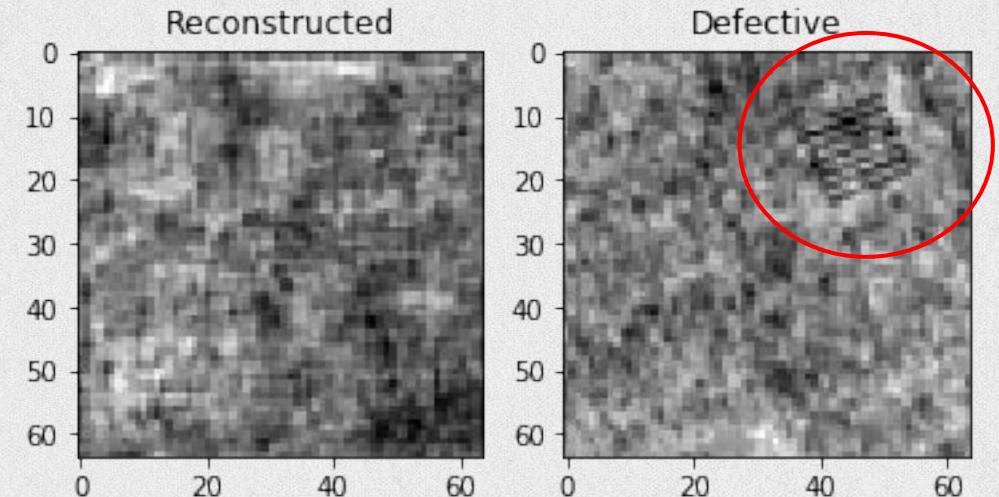
And Johnny ... Does your trick work ?

Healthy Case



Loss: 361.672

Defective Case



2.77x more

Loss: 1002.677

5. Tips n Tricks to effectively train a GAN



5. Tips n Tricks to effectively train a GAN

1. Fractionally Strided Convolution vs Resize + Convolution

```
def deconv2Dlayer(self, x, nb_filters):
    x = Conv2DTranspose(
        filters      = nb_filters,
        kernel_size = (5, 5),
        strides     = (2, 2),
        padding     = 'same')
    ) (x)

    x = LeakyReLU(alpha=0.2) (x)
    x = BatchNormalization(momentum=0.8) (x)

    return x
```

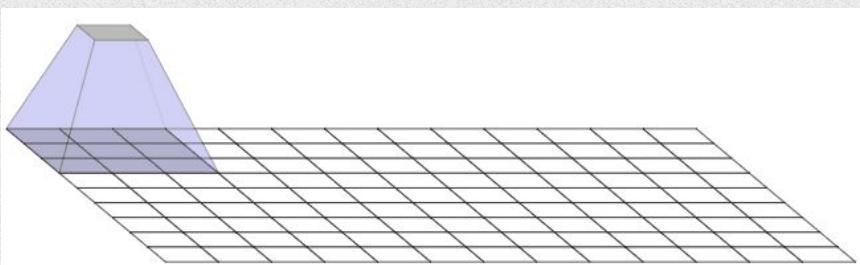
5. Tips n Tricks to effectively train a GAN

1. Fractionally Strided Convolution vs Resize + Convolution

```
def deconv2Dlayer(self, x, nb_filters):  
    x = Conv2DTranspose(  
        filters      = nb_filters,  
        kernel_size = (5, 5),  
        strides     = (2, 2),  
        padding     = 'same'  
    )(x)  
  
    x = LeakyReLU(alpha=0.2)(x)  
    x = BatchNormalization(momentum=0.8)(x)  
  
    return x
```

Fractionally Strided Convolution creates artifacts.

** Inherent to the deconvolution operation
** No “setting” can prevent this



Source: <https://distill.pub/2016/deconv-checkerboard/>

5. Tips n Tricks to effectively train a GAN

1. Fractionally Strided Convolution vs Resize + Convolution

```
def deconv2Dlayer(self, x, nb_filters):
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(
        filters      = nb_filters,
        kernel_size = (5, 5),
        strides     = (1, 1),
        padding     = 'same'
    )(x)

    x = LeakyReLU(alpha=0.2)(x)
    x = BatchNormalization(momentum=0.8)(x)

    return x
```

5. Tips n Tricks to effectively train a GAN

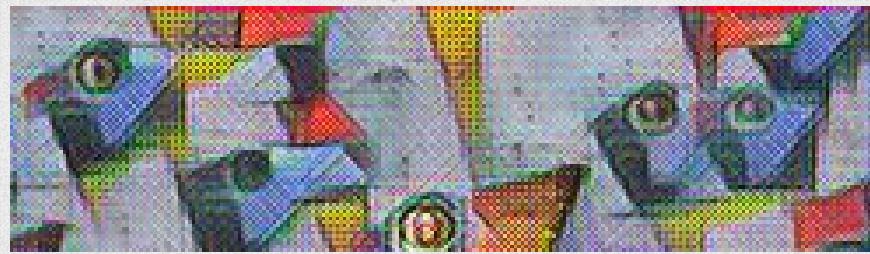
1. Fractionally Strided Convolution vs Resize + Convolution

```
def deconv2Dlayer(self, x, nb_filters):
    x = UpSampling2D((2, 2))(x)
    x = Conv2D(
        filters      = nb_filters,
        kernel_size = (5, 5),
        strides     = (1, 1),
        padding     = 'same'
    )(x)

    x = LeakyReLU(alpha=0.2)(x)
    x = BatchNormalization(momentum=0.8)(x)

    return x
```

With Fractionally Strided Convolutions



With Resize + Convolutions



Source: <https://distill.pub/2016/deconv-checkerboard/>

5. Tips n Tricks to effectively train a GAN

2. Randomly Sampling Z from a non-uniform latent space

The Generator G creates very similar-looking images for close positions in the latent-space:

- ⇒ For certain objectives, we want that a small change in Z lead to a large variation
- ⇒ How To: Train G with Z vectors sampled from a non uniform latent-space
 - ⇒ Sample from a *Truncated Gaussian Distribution* for instance

```
from scipy.stats import truncnorm

def get_truncated_normal(mean=0, std=0.6, low=-1., high=1.):
    # Constrain the sampled data to be in [-1., 1.], centered on 0 with a std of 0.6 (hyperparameter)
    return truncnorm((low - mean) / std, (high - mean) / std, loc=mean, scale=std)

latent_generator = get_truncated_normal()

z = latent_generator.rvs( size=(64, 100) ) # generate a batch of 64 z vectors (64, 100)
```

5. Tips n Tricks to effectively train a GAN

3. Avoid to kill your gradients (sparse) => ReLU & MaxPool

GAN training leads to an instable equilibrium, avoid everything that makes it less stable:

- ⇒ **ReLU** kills the gradients:
- ⇒ **MaxPooling**: the gradient from the next layer is passed back to only the neuron which achieved the max. All other neurons get zero gradient.

Better using LeakyReLU => for G and D.

- ⇒ **Breaking news:** Now implemented in TensorFlow 1.4.0-rc0: `tf.nn.leaky_relu`

5. Tips n Tricks to effectively train a GAN

4. Try everything you can do to make the discriminator's job harder

Forging some data is far harder than discriminating legit from fake data => D trains faster.

- ⇒ Train G twice each time you train D.
- ⇒ Train D only if G accuracy is higher than a threshold (0.3 ? 0.25 ? => HyperParameter)
- ⇒ Use a less efficient optimizer for D (SGD) than for G (Adam).
- ⇒ Noisy Labels for D: Flip labels from time to time for D: Fake is Legit and Legit is fake
- ⇒ Soft Labels for D: Do not use “{0 ,1}” labels but “{0.1, 0.9}”
- ⇒ Add noise for D and decay it over time (Dropout, GaussianNoise, etc.)

THANK YOU

Jonathan DEKHTIAR - @Born2Data
contact@jonathandekhtiar.eu

“Let's see if I can back-propagate
from the questions to the answers.”